

```

import pandas as pd

# Load the dataset (update the path if needed)
df = pd.read_csv('/content/MyProjects/predictive_maintenance.csv')

# Inspect the data
print("Shape of dataset:", df.shape)
print("Columns:", df.columns)
print(df.head())

# Check for missing values
print(df.isnull().sum())

# Check for duplicates
print("Duplicates:", df.duplicated().sum())

```

```

Shape of dataset: (10000, 10)
Columns: Index(['UDI', 'Product ID', 'Type', 'Air temperature [K]', 'Process temperature [K]', 'Rotational speed [rpm]', 'Torque [Nm]', 'Tool wear [min]', 'Target', 'Failure Type'], dtype='object')
   UDI  Product ID  Type  Air temperature [K]  Process temperature [K] \
0     1        M14860    M            298.1                  308.6
1     2        L47181    L            298.2                  308.7
2     3        L47182    L            298.1                  308.5
3     4        L47183    L            298.2                  308.6
4     5        L47184    L            298.2                  308.7

   Rotational speed [rpm]  Torque [Nm]  Tool wear [min]  Target  Failure Type
0                 1551       42.8             0           0      No Failure
1                 1408       46.3             3           0      No Failure
2                 1498       49.4             5           0      No Failure
3                 1433       39.5             7           0      No Failure
4                 1408       40.0             9           0      No Failure
UDI
Product ID
Type
Air temperature [K]
Process temperature [K]
Rotational speed [rpm]
Torque [Nm]
Tool wear [min]
Target
Failure Type
dtype: int64
Duplicates: 0

```

## Data Cleaning & Preparation

```

# Drop duplicates
df = df.drop_duplicates()
print("After dropping duplicates:", df.shape)

```

```
After dropping duplicates: (10000, 10)
```

```
#handle missing values
print(df.isnull().sum())
```

```

UDI          0
Product ID  0
Type         0
Air temperature [K]  0
Process temperature [K]  0
Rotational speed [rpm]  0
Torque [Nm]  0
Tool wear [min]  0
Target        0
Failure Type  0
dtype: int64

```

```

df.columns = df.columns.str.strip().str.lower().str.replace(' ', '_')
print(df.columns)

```

```
Index(['udi', 'product_id', 'type', 'air_temperature_[k]',  
       'process_temperature_[k]', 'rotational_speed_[rpm]', 'torque_[nm]',  
       'tool_wear_[min]', 'target', 'failure_type'],  
      dtype='object')
```

```
#converted to numeric codes for analysis/visualization  
df['product_id_code'] = df['product_id'].map({'L': 1, 'M': 2, 'H': 3})  
df['failure_type_code'] = df['failure_type'].astype('category').cat.codes
```

## Exploratory Data Analysis (EDA)

```
import matplotlib.pyplot as plt  
import seaborn as sns
```


```
# Quick stats of numeric features  
print(df.describe())
```

```
# Count of failures (target = 1 means failure)  
print(df['target'].value_counts())
```


```
      udi  air_temperature_[k]  process_temperature_[k]  \\\ncount  10000.00000    10000.00000    10000.00000  
mean    5000.50000    300.004930   310.005560  
std     2886.89568    2.000259    1.483734  
min     1.00000    295.300000   305.700000  
25%    2500.75000    298.300000   308.800000  
50%    5000.50000    300.100000   310.100000  
75%    7500.25000    301.500000   311.100000  
max    10000.00000    304.500000   313.800000  
  
      rotational_speed_[rpm]  torque_[nm]  tool_wear_[min]  target  \\\ncount    10000.000000  10000.000000  10000.000000  10000.00000  
mean     1538.776100   39.986910   107.951000   0.033900  
std      179.284096   9.968934   63.654147   0.180981  
min     1168.000000   3.800000   0.000000   0.000000  
25%    1423.000000   33.200000   53.000000   0.000000  
50%    1503.000000   40.100000   108.000000   0.000000  
75%    1612.000000   46.800000   162.000000   0.000000  
max    2886.000000   76.600000   253.000000   1.000000  
  
      product_id_code  failure_type_code  
count          0.0        10000.000000  
mean         NaN        1.039000  
std          NaN        0.379069  
min          NaN        0.000000  
25%          NaN        1.000000  
50%          NaN        1.000000  
75%          NaN        1.000000  
max          NaN        5.000000  
target  
0    9661  
1     339  
Name: count, dtype: int64
```

```
# Failure count by type  
plt.figure(figsize=(8,5))  
sns.countplot(x='failure_type', data=df)  
plt.title('Failure Count by Type')  
plt.show()
```

### Failure Count by Type




```
# Failures by Product Variant (product_id)
plt.figure(figsize=(8,5))
sns.countplot(x='product_id', hue='target', data=df)
plt.title('Failures by Product Variant')
plt.show()
```




```
# Tool wear by Product Variant
plt.figure(figsize=(8,5))
sns.boxplot(x='product_id', y='tool_wear_[min]', data=df)
plt.title('Tool Wear by Product Variant')
plt.show()
```

### Tool Wear by Product Variant




```
# Distribution of numeric features
numeric_features = ['air_temperature_[k]', 'process_temperature_[k]',
                     'rotational_speed_[rpm]', 'torque_[nm]', 'tool_wear_[min]']

plt.figure(figsize=(15,10))
for i, col in enumerate(numeric_features):
    plt.subplot(2,3,i+1)
    sns.histplot(df[col], kde=True)
    plt.title(f'{col} Distribution')
plt.tight_layout()
plt.show()
```



```
# Correlation heatmap with target
plt.figure(figsize=(10,6))
sns.heatmap(df[numerical_features + ['target']].corr(), annot=True, cmap='coolwarm')
plt.title('Feature Correlation with Failure')
plt.show()
```



```
# MTBF per Product Variant
mtbf = df.groupby('product_id')['target'].apply(lambda x: 1/(x.mean() + 1e-5))
print("MTBF per Product Variant:\n", mtbf)
```

```
MTBF per Product Variant:
product_id
H29424    1000000.0
H29425    1000000.0
H29432    1000000.0
H29434    1000000.0
H29441    1000000.0
...
M24849    1000000.0
M24851    1000000.0
M24855    1000000.0
M24857    1000000.0
M24859    1000000.0
Name: target, Length: 10000, dtype: float64
```

```
# Average tool wear per Product Variant
avg_tool_wear = df.groupby('product_id')['tool_wear_[min]'].mean()
print("Average Tool Wear per Product Variant:\n", avg_tool_wear)
```

```
Average Tool Wear per Product Variant:
product_id
H29424    24.0
H29425    29.0
H29432    50.0
H29434    58.0
H29441    77.0
...
M24849     0.0
M24851     5.0
M24855    14.0
```

```
M24857    22.0
M24859    30.0
Name: tool_wear_[min], Length: 10000, dtype: float64
```

## Key Metrics & Insights

```
#Encode Target
df['target_code'] = df['target'] # just to be explicit
```

```
#Mean Time Between Failures (MTBF)
mtbf = df.groupby('product_id')['target_code'].apply(lambda x: 1/(x.mean() + 1e-5))
print("MTBF per Product Variant:\n", mtbf)
```

```
MTBF per Product Variant:
product_id
H29424    100000.0
H29425    100000.0
H29432    100000.0
H29434    100000.0
H29441    100000.0
...
M24849    100000.0
M24851    100000.0
M24855    100000.0
M24857    100000.0
M24859    100000.0
Name: target_code, Length: 10000, dtype: float64
```

Interpretation: Higher MTBF → less frequent failures → more reliable machines/products.

```
avg_tool_wear = df.groupby('product_id')['tool_wear_[min]'].mean()
print("Average Tool Wear per Product Variant:\n", avg_tool_wear)
```

```
Average Tool Wear per Product Variant:
product_id
H29424    24.0
H29425    29.0
H29432    50.0
H29434    58.0
H29441    77.0
...
M24849    0.0
M24851    5.0
M24855    14.0
M24857    22.0
M24859    30.0
Name: tool_wear_[min], Length: 10000, dtype: float64
```

Insight: High tool wear variants may need more preventive maintenance.

Approximate Overall Equipment Effectiveness (OEE)

Simplified formula: OEE=Availability×Performance×Quality

We can approximate:

Availability: Proportion of time the machine is running (1 – % downtime). Use tool wear as proxy for downtime.

Performance: Actual output / Ideal output (we can approximate rotational speed vs max).

Quality: 1 – Failure Rate

```
# Availability: 1 - normalized tool wear
availability = 1 - (df['tool_wear_[min]'] / df['tool_wear_[min]'].max())

# Performance: normalized rotational speed
performance = df['rotational_speed_[rpm]'] / df['rotational_speed_[rpm]'].max()

# Quality: 1 - failure rate
quality = 1 - df['target_code']

... . . . --
```

```

# Approximate OEE per row
df['OEE'] = availability * performance * quality

# Average OEE per product variant
oee_per_product = df.groupby('product_id')['OEE'].mean()
print("Average OEE per Product Variant:\n", oee_per_product)

```

```

Average OEE per Product Variant:
product_id
H29424    0.558890
H29425    0.436552
H29432    0.363097
H29434    0.367215
H29441    0.436530
...
M24849    0.546431
M24851    0.620545
M24855    0.525032
M24857    0.520428
M24859    0.458120
Name: OEE, Length: 10000, dtype: float64

```

Interpretation: Higher OEE → better overall machine effectiveness.

```

# Top 5 machines/products with highest failure rate
high_risk_products = df.groupby('product_id')['target_code'].mean().sort_values(ascending=False).head(5)
print("High-Risk Product Variants:\n", high_risk_products)

```

```

High-Risk Product Variants:
product_id
L51335    1.0
M21538    1.0
L49255    1.0
L51337    1.0
M19416    1.0
Name: target_code, dtype: float64

```

## Reporting

```

# MTBF per Product
mtbf_table = mtbf.reset_index().rename(columns={'target_code':'MTBF'})
print(mtbf_table)

# Average Tool Wear per Product
tool_wear_table = avg_tool_wear.reset_index().rename(columns={'tool_wear_[min]':'Avg_Tool_Wear'})
print(tool_wear_table)

# OEE per Product
oee_table = oee_per_product.reset_index().rename(columns={'OEE':'Avg_OEE'})
print(oee_table)

# High-risk Products
high_risk_table = high_risk_products.reset_index().rename(columns={'target_code':'Failure_Rate'})
print(high_risk_table)

```

	product_id	MTBF
0	H29424	100000.0
1	H29425	100000.0
2	H29432	100000.0
3	H29434	100000.0
4	H29441	100000.0
...	...	...
9995	M24849	100000.0
9996	M24851	100000.0
9997	M24855	100000.0
9998	M24857	100000.0
9999	M24859	100000.0

[10000 rows x 2 columns]

	product_id	Avg_Tool_Wear
0	H29424	24.0
1	H29425	29.0
2	H29432	50.0
3	H29434	58.0

```
4      H29441    77.0
...
9995   M24849    0.0
9996   M24851    5.0
9997   M24855   14.0
9998   M24857   22.0
9999   M24859   30.0
```

```
[10000 rows x 2 columns]
   product_id  Avg_OEE
0      H29424  0.558890
1      H29425  0.436552
2      H29432  0.363097
3      H29434  0.367215
4      H29441  0.436530
...
9995   M24849  0.546431
9996   M24851  0.620545
9997   M24855  0.525032
9998   M24857  0.520428
9999   M24859  0.458120
```

```
[10000 rows x 2 columns]
   product_id  Failure_Rate
0      L51335     1.0
1      M21538     1.0
2      L49255     1.0
3      L51337     1.0
4      M19416     1.0
```

```
mtbf_table.to_csv('MTBF_per_Product.csv', index=False)
tool_wear_table.to_csv('ToolWear_per_Product.csv', index=False)
oee_table.to_csv('OEE_per_Product.csv', index=False)
```